

Programming in Modern C++: Assignment Week 12

Total Marks : 27

Partha Pratim Das
Department of Computer Science and Engineering
Indian Institute of Technology Kharagpur, Kharagpur – 721302
partha.p.das@gmail.com

April 16, 2025

Question 1

Consider the program (in C++11) given below.

[MSQ, Marks 2]

```
#include <iostream>
#include <thread>
#include <list>
using namespace std;

struct Product {
    list<int>& iLst;
    int* result;
    Product(list<int>& il, int* res) : iLst{il}, result{res} { }
    void operator()() {
        for(auto& i : iLst) {
            *result *= i;
            i = 0;
        }
    }
};

void print(list<int> iLst) {
    for(auto i : iLst)
        std::cout << i << " ";
}

int main() {
    list<int> il { 10, 20, 30 };
    int res = 1;
    std::thread t { _____ };    //LINE-1
    t.join();
    print(il);
    std::cout << res;
    return 0;
}
```

Fill in the blank at LINE-1 with appropriate option(s) such that the output becomes 0 0 0 6000.

- a) `Product(&li, &res)`
- b) `Product(li, &res)`
- c) `Product(std::ref(li), &res)`
- d) `std::bind(Product, std::ref(li), &res)`

Answer: b), c)

Explanation:

The constructor of `Product` receives `li` as pass-by-reference and `res` as pass-by-address.

At option a), the actual parameter `&li` is an address. Therefore, it is a wrong option.

At option b), the actual parameter `li` is a correct way to perform pass-by-reference (or pass-by-value) and `&res` is passed-by-address. Therefore, it is a correct option.

At option c), the actual parameter `std::ref(li)` is a pass-by-reference explicitly and `&res` is passed-by-address. Therefore, it is a correct option.

At option d), the actual parameter `Product` of `std::bind` function is a call to the default constructor of functor `Product`, which is not provided. Therefore, it is a wrong option.

Question 2

Consider the program (in C++11) given below.

[MCQ, Marks 2]

```
#include <iostream>
#include <functional>
using namespace std;

int print(int n1, int n2, int n3, const int& n4) {
    return (n1 - (n2 + n3 + n4));
}

int main() {
    using namespace std::placeholders;
    int val = 10;
    auto f = std::bind(print, _2, val, _1, std::cref(val));
    val = -10;
    std::cout << f(100, 20, -100);
    return 0;
}
```

What will be the output?

- a) -100
- b) -80
- c) 80
- d) 120

Answer: b)

Explanation:

The call `f(100, 20, -100)` results in binding 100 to `_1`, 20 to `_2`, and -100 remains unused. The formal arguments to the call to `print` are as `print(20, 10, 100, -10)`, which is evaluated as -80.

Question 3

Consider the following class (int C++11).

[MCQ, Marks 2]

```
#include <iostream>
#include <thread>
#include <mutex>
#include <functional>

std::mutex val_mutex;
int val = 1;

void stepUp(const int& steps){
    std::unique_lock<std::mutex> lck(val_mutex);
    int i = val;
    i += steps;
    val = i;
}

int main() {
    std::thread t1 { &stepUp, 1 };
    std::thread t2 { &stepUp, 2 };
    std::thread t3 { &stepUp, 3 };
    t1.join();
    t2.join();
    t3.join();
    std::cout << val;
    return 0;
}
```

What will be the output?

- a) 1
- b) 6
- c) 7
- d) Unpredictable

Answer: c)

Explanation:

Since within function `stepUp` all the modifications to `val` takes place in a mutual exclusive way, each thread will update the value of `val` and the result will be 7. However, the order in which order threads update `val` cannot be predicted.

Question 4

Consider the following code segment (in C++11).

[MCQ, Marks 2]

```
#include <iostream>
#include <thread>
#include <mutex>
#include <functional>

std::mutex m;
thread_local int val = 1;
void stepUp(const int& steps){
    m.lock();
    int i = val;
    i += steps;
    val = i;
    m.unlock();
}

int main() {
    val = 0;
    std::thread t1 { std::bind(stepUp, 1) };
    std::thread t2 { std::bind(stepUp, 2) };
    std::thread t3 { std::bind(stepUp, 3) };
    t1.join();
    t2.join();
    t3.join();
    std::cout << val;
    return 0;
}
```

What will be the output?

- a) 0
- b) 1
- c) 6
- d) 7

Answer: a)

Explanation:

Since the definition `thread_local int val = 1;` makes `val` local to each of the threads, each thread increments its local copy of `val`. Thus, output becomes 100 in `main`.

Question 5

Consider the following code segment (in C++11).

[MSQ, Marks 2]

```
#include <iostream>
#include <thread>
#include <mutex>

struct Printer{
    int pCounter = 0;
};

struct Scanner{
    int sCounter = 0;
};

std::mutex printer_mutex;
std::mutex scanner_mutex;
Printer printer;
Scanner scanner;

void requestCopy1(int n) {
    std::unique_lock<std::mutex> lck1(printer_mutex);
    std::unique_lock<std::mutex> lck2(scanner_mutex);
    printer.pCounter += n;
    scanner.sCounter += n;
    std::cout << printer.pCounter << ", " << scanner.sCounter << std::endl;
}

void requestCopy2(int n) {
    std::unique_lock<std::mutex> lck1(scanner_mutex);
    std::unique_lock<std::mutex> lck2(printer_mutex);
    printer.pCounter += n;
    scanner.sCounter += n;
    std::cout << printer.pCounter << ", " << scanner.sCounter << std::endl;
}

int main(){
    std::thread t1(requestCopy1, 10);
    std::thread t2(requestCopy2, 20);
    t1.join();
    t2.join();
    return 0;
}
```

Which of the following cannot be true about the output of the above program?

- a) It generates output as:
10, 10
30, 30
- b) It generates output as:
20, 20
30, 30

c) It generates output as:

20, 20

10, 10

d) It results in deadlock

Answer: a), b)

Explanation:

Since the code in `requestCopy1` and `requestCopy2` execute in a mutual exclusive manner, the output can be:

10, 10

30, 30

or

20, 20

30, 30

However, it may happen that `t1` holds lock on `printer_mutex`, and `t2` holds lock on `scanner_mutex`. Then, `t1` request to lock on `scanner_mutex`, and `t2` requests lock on `scanner_mutex`. It results in a deadlock.

Question 6

Consider the following code segment (in C++11).

[MSQ, Marks 2]

```
#include <iostream>
#include <vector>
#include <thread>

void add_to_vec(std::vector<int>& iVec){
    for(int i = 3; i < 5; i++)
        iVec.push_back(i);
}

int main(){
    std::vector<int> vc;
    for(int i = 1; i < 3; i++)
        vc.push_back(i);
    std::thread t1 {add_to_vec, std::ref(vc)};
    for(auto i : vc)
        std::cout << i << " ";
    t1.join();
    return 0;
}
```

Which of the following **cannot** be valid output?

- a) 1 2
- b) 1 2 3 4
- c) 1 3 2 4
- d) 3 4 1 2

Answer: c), d)

Explanation:

The `main` thread adds 1 and 2 to vector `vc` before the thread `t` got initiated. The thread `t` and the `for` loop printing `vc` work in an interleaved manner. Thus, the options c) and d) are not possible.

Question 7

Consider the following program (in C++11) given below.

[MCQ, Marks 2]

```
#include <iostream>

template <class T>
class SmartPtr {
    private:
        T* _pointee;
    public:
        explicit SmartPtr(T* p = nullptr) { _pointee = p; }
        ~SmartPtr() { delete _pointee; }
        ----- { return *_pointee; }    //LINE-1
};

int main() {
    const SmartPtr<char> sp(new char);
    *sp = 'X';
    std::cout << *sp;
    return 0;
}
```

Identify the correct function header at LINE-1 so that it prints **X**.

- a) T operator* () const
- b) T& operator* () const
- c) T* operator->() const
- d) SmartPtr& operator=

Answer: b)

Explanation:

The statements `*sp = 'X';` as well as `std::cout << *sp;` require dereferencing operator to be overridden. Thus, option b) is the correct option.

Question 8

Consider the following code segment (in C++11).

[MCQ, Marks 2]

```
#include<iostream>
#include <memory>

void update_share_ptr(const std::shared_ptr<int>& sp){
    std::shared_ptr<int> lp = sp;
    std::cout << "RC = " << lp.use_count() << " ";
}

int main(){
    std::shared_ptr<int> spA(new int(10));
    {
        std::shared_ptr<int> spB(spA);
        std::cout << "RC = " << spA.use_count() << " ";
    }
    std::shared_ptr<int> spC = spA;
    std::cout << "RC = " << spA.use_count() << " ";
    update_share_ptr(spA);
    std::cout << "RC = " << spA.use_count() << " ";
    spC.reset(new int(20));
    std::cout << "RC = " << spA.use_count() << " ";
    return 0;
}
```

What will be the output?

- a) RC = 1 RC = 1 RC = 2 RC = 2 RC = 1
- b) RC = 2 RC = 2 RC = 4 RC = 2 RC = 1
- c) RC = 2 RC = 2 RC = 4 RC = 3 RC = 2
- d) RC = 2 RC = 2 RC = 3 RC = 2 RC = 1

Answer: d)

Explanation:

The code is explained in the comment:

```
#include<iostream>
#include <memory>

void update_share_ptr(const std::shared_ptr<int>& sp){ //RC = 2 since pass-by-reference
    std::shared_ptr<int> lp = sp; //RC = 3
    std::cout << "RC = " << lp.use_count() << " ";
}

int main(){
    std::shared_ptr<int> spA(new int(10)); //RC = 1
    {
        std::shared_ptr<int> spB(spA); //RC = 2
        std::cout << "RC = " << spA.use_count() << " ";
    }
    std::shared_ptr<int> spC = spA; //RC = 2 since spB got deleted
    std::cout << "RC = " << spA.use_count() << " ";
}
```

```

    update_share_ptr(spA);                //call to update_share_ptr
    std::cout << "RC = " << spA.use_count() << " ";
    spC.reset(new int(20));                //RC = 1 since local lp got deleted
    std::cout << "RC = " << spA.use_count() << " ";
    return 0;
}

```

Question 9

Consider the following code segment (in C++).

[MCQ, Marks 2]

```
#include <iostream>
#include <memory>

void update_share_ptr1(std::shared_ptr<int>& sp){
    sp.reset(new int(3));
}

void update_share_ptr2(std::shared_ptr<int> sp){
    sp.reset(new int(4));
}

int main() {
    auto sp1 = std::make_shared<int>(1);
    auto sp2 = sp1;
    std::weak_ptr<int> wp1 = sp1;
    std::weak_ptr<int> wp2 = sp2;
    update_share_ptr1(sp1);
    update_share_ptr2(sp2);

    if(auto p = wp1.lock())
        std::cout << *p << std::endl;
    else
        std::cout << "wp1 is expired" << std::endl;
    update_share_ptr1(sp2);
    update_share_ptr2(sp1);

    if(auto p = wp2.lock())
        std::cout << *p << std::endl;
    else
        std::cout << "wp2 is expired" << std::endl;
    return 0;
}
```

What will be the output?

- a) 1
2
- b) wp1 is expired
1
- c) 1
wp2 is expired
- d) wp1 is expired
wp2 is expired

Answer: c)

Explanation:

Since function `update_share_ptr1` uses pass-by-reference as parameter, it resets the actual shared pointer passed as parameter.

Since function `update_share_ptr2` uses pass-by-value as parameter, it does not affect the actual shared pointer passed as parameter.

The statement `auto sp2 = sp1;` sets `RC = 2` for `sp1` as well as `sp2`. Weak pointers `wp1` and `wp2` refers to `sp1` and `sp2`.

The statements `update_share_ptr1(sp1);` and `update_share_ptr2(sp2);` results in `RC = 1` (since the second call does not affect the share pointers in `main`). Thus, it prints 1 as the first output.

Next, the statements `update_share_ptr2(sp1);` and `update_share_ptr1(sp2);` results in `RC = 0`. Thus, it prints `wp2 is expired`.

Programming Questions

Question 1

Consider the program below (in C++11).

- Fill in the blank at LINE-1 and LINE-2 with appropriate definitions of `std::promise p1` and `std::future f1`.
- Fill in the blank at LINE-3 and LINE-4 with appropriate definitions of `std::promise p2` and `std::future f2`.
- Fill in the blank at LINE-5 to define a `std::thread` object by passing appropriate `std::future` and `std::promise` objects.

The program must satisfy the given test cases.

Marks: 3

```
#include <iostream>
#include <functional>
#include <thread>
#include <future>
#include <vector>

void product (std::promise<long>& p, std::future<std::vector<int>>& f) {
    std::vector<int> v = f.get();
    long result = 1;
    for(int i : v)
        result *= (long)i;
    p.set_value(result);
}

int main () {
    -----; //LINE-1
    -----; //LINE-2

    -----; //LINE-3
    -----; //LINE-4

    -----; //LINE-5

    int n, a;
    std::cin >> n;
    std::vector<int> v;
    for(int i = 0; i < n; i++){
        std::cin >> a;
        v.push_back(a);
    }

    p1.set_value (v);
    std::cout << f2.get();
    th.join();
    return 0;
}
```

Public 1

Input: 3 10 20 30

Output: 6000

Public 2

Input: 4 10 -10 20 20

Output: -40000

Private

Input: 5 2 -3 -4 5 6

Output: 720

Answer:

```
LINE-1:  std::promise<std::vector<int>> p1
LINE-2:  std::future<std::vector<int>> f1 = p1.get_future()
LINE-3:  std::promise<long> p2
LINE-4:  std::future<long> f2 = p2.get_future()
LINE-5:  std::thread th (product, std::ref(p2), std::ref(f1))
```

Explanation:

At LINE-1 and LINE-2, the `std::promise p1` and `std::future f1` can be created as follows:

```
LINE-1:  std::promise<std::vector<int>> p1
LINE-2:  std::future<std::vector<int>> f1 = p1.get_future()
```

At LINE-3 and LINE-4, the `std::promise p2` and `std::future f2` can be created as follows:

```
LINE-3:  std::promise<long> p2;
LINE-4:  std::future<long> f2 = p2.get_future()
```

The function `product` to get the values from `main` using `std::promise p1`, whereas it sets the value to `std::future f2` so that it can be accessed by `main`. Therefore, the thread can be defined as follows

```
LINE-5:  std::thread th (product, std::ref(p2), std::ref(f1))
```

Question 2

Consider the following program (in C++14).

- Fill in the blank at LINE-1 with an appropriate definition of mutex `mq_mutex`.
- Fill in the blanks at LINE-2 and LINE-3 with appropriate statements such that all the modification to `total_elems` happens in a mutual exclusive manner.

The program must satisfy the sample input and output.

Marks: 3

```
#include <iostream>
#include <vector>
#include <thread>
#include <functional>
#include <chrono>
#include <mutex>
-----;    //LINE-1

class MessageQueue{
private:
    int total_elems;
    int n;
public:
    MessageQueue() : total_elems(0) {}
    void enqueue(int num_elems){
        -----;    //LINE-2
        n = num_elems;
        int delay = (int)((double)std::rand() / (double)(RAND_MAX)* 10);
        std::this_thread::sleep_for(std::chrono::milliseconds(delay));
        total_elems += n;
    }
    void dequeue(int num_elems){
        -----;    //LINE-3
        n = num_elems;
        int delay = (int)((double)std::rand() / (double)(RAND_MAX)* 10);
        std::this_thread::sleep_for(std::chrono::milliseconds(delay));
        total_elems -= n;
    }
    int getElementsCount(){
        return total_elems;
    }
};

void sender(MessageQueue& mq, int n){
    for(int i = 0; i < n; i++){
        int delay = (int)((double)std::rand() / (double)(RAND_MAX)* 10);
        std::this_thread::sleep_for(std::chrono::milliseconds(delay));
        mq.enqueue(i + 1);
    }
}

void receiver(MessageQueue& mq, int n){
    for(int i = n - 1; i >= 0; i--){
```



```

        int delay = (int)((double)std::rand() / (double)(RAND_MAX)* 30);
        std::this_thread::sleep_for(std::chrono::milliseconds(delay));
        mq.dequeue(i + 1);
    }
}

int main(){
    int n, m;
    std::cin >> n;
    std::cin >> m;
    MessageQueue mq;
    std::thread t1{ std::bind(sender, std::ref(mq), n) };
    std::thread t2{ std::bind(receiver, std::ref(mq), m) };

    t1.join();
    t2.join();
    std::cout << mq.getElemsCount();
    return 0;
}

```

Public 1

Input: 10 10

Output: 0

Public 2

Input: 20 10

Output: 155

Public 3

Input: 10 20

Output: -155

Private

Input: 20 20

Output: 0

Answer:

LINE-1: `std::mutex mq_mutex`

LINE-2: `std::unique_lock<std::mutex> lck(mq_mutex)`

LINE-3: `std::unique_lock<std::mutex> lck(mq_mutex)`

Explanation:

At LINE-1, mutex can be defined as:

`std::mutex mq_mutex`

At LINE-2 and LINE-3 we can use the following statement:

`std::unique_lock<std::mutex> lck(mq_mutex)` such that any update to the variable `total_elem` would be done in a critical section.

Question 3

Consider the following program (in C++11).

- Fill in the blank at LINE-1 with appropriate header to overload function operator.
- Fill the blank at LINE-2 to invoke the functor `NthPrime()` asynchronously.
- Fill the blank at LINE-3 to receive the output from functor `NthPrime()`.

The program must satisfy the sample input and output.

Marks: 3

```
#include <iostream>
#include <future>
#include <cmath>
using namespace std;

struct NthPrime{
    const long long n;
    NthPrime(const long long& n) : n( n) { }
    bool isPrime(long long n){
        if (n <= 1)
            return false;
        for (int i = 2; i <= sqrt(n); i++)
            if (n % i == 0)
                return false;

        return true;
    }
    ----- { //LINE-1
        long long num = 2;
        for(int i = 0; i < n; num++) {
            if (isPrime(num)) {
                ++i;
            }
        }
        return num-1;
    }
};

long long findNthPrime(int n){
    ----- ; //LINE-2
    return ----- ; //LINE-3
}

int main() {
    int n;
    cin >> n;
    cout << findNthPrime(n) << endl;
    return 0;
}
```

Public 1

Input: 150

Output: 863

Public 2

Input: 200

Output: 1223

Private

Input: 1000

Output: 7919

Answer:

LINE-1: `long long operator()()`

LINE-2: `auto a = async(NthPrime(n))`

LINE-3: `a.get()`

Explanation:

The function header at LINE-1 to overload the function operator can be: `long long operator()()`

At LINE-2 the asynchronous call to the functor `NthPrime()` can be made as:

`auto a = async(NthPrime(n))`

At LINE-3 can use the statement:

`a.get()`

to receive the result of functor `NthPrime()`.